

# Designing Data Intensive Applications

## Chapter 1

# Reliability

- Tolerating hardware and software faults
- Human error

# Scalability

- Measuring load & Performance
- Latency percentiles, throughput

# Maintainability

- Operability, simplicity, & evolvability

# Common requirements of data-intensive applications

- Store data (databases)
- Cache expensive operations (memory stores)
- Allow search (indexes)
- Asynchronously sending messages (stream processing)
- Crunch data (batch processing)

# Why data systems?

- Tools for data storage and processing are specialized for specific uses.
- No single tool can satisfy all requirements.
- API hides implementation details.

# Concerns of data-intensive applications

- Data remains correct and complete
- Provide consistently good performance to clients
- Handle increased load
- User friendly API

# Reliability

- The application performs the function that the user expected.
- It can tolerate the user making mistakes or using the software in unexpected ways.
- Its performance is good enough for the required use case, under the expected load and data volume.
- The system prevents any unauthorized access and abuse.



## Some quick definitions

Fault is one component of the system deviating from its spec.

Failure is when the system as a whole stops providing the required service to the user.

# Chaos Monkey

- Increasing the rate of faults can improve fault-tolerance.
- Google's enforced SLA when they exceed uptime.

# Hardware faults

"Hard disks are reported as having a mean time to failure (MTTF) of about 10 to 50 years. Thus, on a storage cluster with 10,000 disks, we should expect on average one disk to die per day."

**Hardware fault solution?**

# Redundancy

- RAID
- Hot swap CPU
- Hot swap power supply
- Multiple NICs

# Software Errors

- Bugs
- Resource hogs
- Bottlenecks (service is unresponsive or slow)
- Cascading failures (stampeding herd's etc.)

Software Errors are a surprise. Mitigate them by...

- Testing
- Process isolation
- Let it crash!
- Measuring/monitoring

# Human Errors

- Design systems in a way that minimizes opportunities for error. (duh?)
- Decouple environments.
- Test Test Test (Balance is hard)
- Enable easy recovery (quick config rollbacks)
- Detailed telemetry
- Runbooks



# Scalability is tradeoffs

How can we add resources to handle load?

# Load

- Systems have different definitions of load
- rps
- DB reads/writes
- Cache hit rate

# Twitter example!

Post tweet - 4.6k rps avg 12k rps max

Home timeline - 300k rps

SLO tweets appear within 5s.

Twitter Option 1

Twitter Option 2 (fanout)

Twitter Solution

Twitter Solution

Throughput (performance) vs Response Time (responsiveness)



# Five 9's

Uptime calculator

# Amazon cares about 99.9 percentile 1 in 1,000

- Users with the slowest requests have more data, usually mvps
- 99.99 1 in 10,000 was too expensive
- SLOs service level objectives vs SLAs service level agreements
- The slowest backend request is your response time.

# Handling Load changes

10,000 -> 100,000 -> 1,000,000 -> 10,000,000

# Scaling

Vertical vs Horizontal scaling

Elastic systems (autoscaling)

# Maintainability

- Operability - easy to keep the system running
- Simplicity - easy for new engineers to understand the system
- Evolvability - easy to make changes to the system

# Operability

- Monitoring health
- Root causing issues
- Patching
- Understanding system-system side effects
- capacity planning
- Establishing best practices for tools and config
- Maintenance and migration
- Maintaining security
- Predictable operations (deploys)
- Preserve knowledge (Documentation)

# Simplicity

Simplicity is not a lack of functionality

# Symptoms of Complexity

- Large number of states
- tight coupling of modules
- tangled dependencies
- inconsistent naming and terminology
- perf optimization hacks
- special-case bug fixes



# Nothing is Something (my favorite talk on abstractions)

<https://youtu.be/OMPfEXIITVE>

**Evolvability: Making Change Easy**

